

SDS 383D: Exercise 2

Diego Garcia-Olano

February 8, 2017

Problem 1. Bayes and the Gaussian linear model

A simple Gaussian location model

Take a simple Gaussian model with unknown mean and variance:

$$(y_i | \theta, \sigma^2) \sim N(\theta, \sigma^2), i = 1, \dots, n. \quad (1)$$

Let y be the vector of observations $y = (y_1, \dots, y_n)^T$.

Suppose we place conjugate normal and inverse-gamma priors on θ and σ^2 , respectively:

$$p(\theta | \sigma^2) \sim N(\mu, \tau^2 \sigma^2)$$

$$\sigma^2 \sim \text{InvGamma}\left(\frac{d}{2}, \frac{\eta}{2}\right)$$

where $\mu, \tau > 0$, $d > 0$ and $\eta > 0$ are fixed scalar hyperparameters.

*Note a crucial choice here: the error variance σ^2 appears in the prior for θ .

This affects the interpretation of the hyperparameter τ ,

which is not the prior variance of θ , but rather the prior signal-to-noise ratio.

This is pretty common thing to do in setting up priors for location parameters:

to *scale the prior by the error variance*. There are a few good reasons to do this,

but historically the primary one has been analytical convenience (as you'll now see).

Here's a sensible way to interpret each of these four parameters:

- μ is a prior guess for θ .
- τ is a prior signal-to-noise ratio
- that is, how disperse your prior is for θ , relative to the error standard deviation σ .
- d is like a "prior sample size" for the error variance σ^2 .
- η is like a "prior sum of squares" for the error variance σ^2 .
More transparently, η/d is like a "**prior guess**" for the error variance σ^2 . It's not exactly the prior mean for σ^2 , but it's close to the prior mean as d gets larger, since the inverse-gamma(a,b) prior has expected value

$$E(\sigma^2) = \frac{b}{a-1} = \frac{\eta/2}{d/2-1} = \frac{\eta}{d-2}$$

if d is large. This expression is only valid if $d > 2$.

What is meant by "prior sample size" (d) and "prior sum of squares" (η)?

Remember that **conjugate priors always resemble the likelihood functions** that they're intended to play nicely with. The two relevant quantities in the likelihood function for σ^2 are (i) the sample

size and (ii) the sums of squares. The prior here is designed to mimic the likelihood function for σ^2 that you'd get if you had a previous data set with sample size d and sums of squares η .

Precisions are easier than variances. It's perfectly fine to work with this form of the prior, and it's easier to interpret this way. But it turns out that we can make the algebra a bit cleaner by working with the precisions: $\omega = \frac{1}{\sigma^2}$ and $\kappa = \frac{1}{\tau^2}$ instead.

$$p(\theta|\omega) \sim N(\mu, (\omega\kappa)^{-1})$$

$$\omega \sim \text{Gamma}(\frac{d}{2}, \frac{\eta}{2})$$

This means that the joint prior for (θ, ω) has the form:

$$p(\theta, \omega) \propto \omega^{\frac{d+1}{2}-1} \cdot \exp\left(-\omega \frac{\kappa(\theta - \mu)^2}{2}\right) \cdot \exp\left(-\omega \cdot \frac{\eta}{2}\right) \quad (2)$$

This is often called the *normal/gamma* prior for (θ, ω) with parameters (μ, κ, d, η) , and its equivalent to a normal/inverse-gamma prior for (θ, σ^2) .

The interpretation of κ is like a *prior sample size* for the mean θ

Note: you can obviously write this joint density for $p(\theta|\omega)$ in a way that combines the exponential terms, but this way keeps the bit involving θ separate, so that you can recognize the normal kernel. The term "kernel" is heavily overloaded in statistics so see [https://en.wikipedia.org/wiki/Kernel_\(statistics\)#In_Bayesian_statistics](https://en.wikipedia.org/wiki/Kernel_(statistics)#In_Bayesian_statistics).

- (A) By construction, we know that the marginal prior distribution $p(\theta)$ is a gamma mixture of normals. Show that this takes the form of a centered, scaled t distribution:

$$p(\theta) \propto \left(1 + \frac{1}{v} \cdot \frac{(x - m)^2}{s^2}\right)^{-\frac{v+1}{2}}$$

with center m , scale s , and degrees of freedom v ,

where you fill in the blank for m , s^2 , and v in terms of the four parameters of the normal-gamma family. * you did a problem like this in exercises 1!

By definition, the marginal of θ is the integral of the joint distribution $p(\theta, \omega)$ with respect to ω (ie, we integrate out ω) so

$$\begin{aligned} p(\theta) &= \int p(\theta, \omega) d\omega \\ &\propto \int \omega^{\frac{d+1}{2}-1} \cdot \exp\left(-\omega \frac{\kappa(\theta - \mu)^2}{2}\right) \cdot \exp\left(-\omega \cdot \frac{\eta}{2}\right) \\ &= \int \omega^{\frac{d+1}{2}-1} \cdot \exp\left(\frac{\omega \kappa(\theta - \mu)^2 + -\omega \eta}{2}\right) \\ &= \int \omega^{(\frac{d+1}{2}-1)} \cdot \exp\left(-\omega \left(\frac{\kappa(\theta - \mu)^2 + \eta}{2}\right)\right) \end{aligned} \quad (3)$$

This is the kernel for gamma(a,b) for $a = \frac{d+1}{2}$ and $b = \frac{\kappa(\theta-\mu)^2 + \eta}{2}$

As a probability distribution, it must integrate to $1/c$ where c is a constant.

In the case for a gamma distribution the leading constant for the pdf is

$\frac{b^a}{\Gamma(a)} = c$ thus $1/c = \frac{\Gamma(a)}{b^a}$ so we can rewrite (1) as:

$$\begin{aligned}
 &= \Gamma\left(\frac{d+1}{2}\right) \left(\frac{\kappa(\theta-\mu)^2 + \eta}{2}\right)^{-\frac{d+1}{2}} \\
 &\propto \left(\frac{\kappa(\theta-\mu)^2 + \eta}{2}\right)^{-\frac{d+1}{2}} \\
 &= \left(\frac{\eta}{2} + \frac{\kappa(\theta-\mu)^2}{2}\right)^{-\frac{d+1}{2}} \\
 &= \left(\frac{\eta}{2} \cdot \left(1 + \frac{\kappa(\theta-\mu)^2}{\eta}\right)\right)^{-\frac{d+1}{2}} \\
 &= \frac{\eta^{-\frac{d+1}{2}}}{2} \cdot \left(1 + \frac{\kappa(\theta-\mu)^2}{\eta}\right)^{-\frac{d+1}{2}}
 \end{aligned} \tag{4}$$

Since we only care about the value of the equation with respect to θ , we can treat the first part of the equation as constant

$$\begin{aligned}
 &= \left(1 + \frac{\kappa(\theta-\mu)^2}{\eta}\right)^{-\frac{d+1}{2}} \\
 &= \left(1 + \frac{d}{d} \cdot \frac{\kappa(\theta-\mu)^2}{\eta}\right)^{-\frac{d+1}{2}} \\
 &= \left(1 + \frac{1}{d} \cdot \frac{(\theta-\mu)^2}{\eta d/\kappa}\right)^{-\frac{d+1}{2}}
 \end{aligned} \tag{5}$$

This is close to the form of our solution,

$$p(\theta) \propto \left(1 + \frac{1}{v} \cdot \frac{(x-m)^2}{s^2}\right)^{-\frac{v+1}{2}}$$

Thus set $v = d$, $m = \mu$ and set $s^2 = \eta d/\kappa$ to get the centered, scaled t-student form.

(B) Assume the normal sampling model in Eq 1 and the normal-gamma prior in Eq 2.

Calculate joint posterior density $p(\theta, \omega | \mathbf{y})$, up to constant factors not depending on ω or θ .

Show that this is also a normal/gamma prior in the same form as above:

$$p(\theta, \omega | \mathbf{y}) \propto \omega^{(d^*+1)/2-1} \exp\left\{-\omega \cdot \frac{\kappa^*(\theta-\mu^*)^2}{2}\right\} \exp\left\{-\omega \cdot \frac{\eta^*}{2}\right\}$$

We have normal sampling model (eq 1): $(\mathbf{y}|\theta, \sigma^2) \sim N(\theta, \sigma^2)$

and also normal-gamma prior (eq 2) : $p(\theta, \omega) \propto \omega^{\frac{d+1}{2}-1} \cdot \exp\left(-\omega \frac{\kappa(\theta-\mu)^2}{2}\right) \cdot \exp\left(-\omega \cdot \frac{\eta}{2}\right)$

Calculate the joint posterior density,

$$\begin{aligned}
 p(\theta, \omega | \mathbf{y}) &\propto p(\mathbf{y} | \theta, \omega) p(\theta, \omega) \\
 &= \omega^{n/2} \exp\left\{-\omega \cdot \left(\frac{S_y + n(\bar{y} - \theta)^2}{2}\right)\right\} \cdot \omega^{\frac{d+1}{2}-1} \cdot \exp\left(-\omega \frac{\kappa(\theta - \mu)^2}{2}\right) \cdot \exp\left(-\omega \cdot \frac{\eta}{2}\right) \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\omega \left(\frac{S_y + n(\bar{y} - \theta)^2 + \kappa(\theta - \mu)^2 + \eta}{2}\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}(S_y + n\bar{y}^2 + n\theta^2 - 2n\bar{y}\theta + \kappa\theta^2 + \kappa\mu^2 - 2\kappa\theta\mu + \eta)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}((n + \kappa)\theta^2 - 2(n\bar{y} + \kappa\mu)\theta + (S_y + n\bar{y}^2 + \kappa\mu^2 + \eta))\right\}
 \end{aligned} \tag{6}$$

The term in 2nd part that is multiplied by $-\frac{\omega}{2}$ is of the form $ax^2 - 2bx + c$ so we can complete the square as such:

$$\begin{aligned}
 ax^2 - 2bx + c &= a\left[x^2 - 2\left(\frac{b}{a}\right)x + \frac{c}{a}\right] \\
 &= a\left[x^2 - 2\left(\frac{b}{a}\right)x + \left(\frac{b}{a}\right)^2 - \left(\frac{b}{a}\right)^2 + \frac{c}{a}\right] \\
 &= a\left[\left(x - \frac{b}{a}\right)^2 - \left(\frac{b}{a}\right)^2 + \frac{c}{a}\right] \\
 &= a\left(x - \frac{b}{a}\right)^2 - \frac{b^2}{a} + c
 \end{aligned} \tag{7}$$

Now plugging in for $a = n + \kappa$, $x = \theta$, $b = n\bar{y} + \kappa\mu$, and c = the final term in (6), we get:

$$\begin{aligned}
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 - \frac{(n\bar{y} + \kappa\mu)^2}{n + \kappa} + (S_y + n\bar{y}^2 + \kappa\mu^2 + \eta)\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 - \frac{(n^2\bar{y}^2 + \kappa^2\mu^2 + 2n\bar{y}\kappa\mu)}{n + \kappa} + (n\bar{y}^2 + \kappa\mu^2) + S_y + \eta\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 - \frac{(n^2\bar{y}^2 + \kappa^2\mu^2 + 2n\bar{y}\kappa\mu)}{n + \kappa} + \frac{n + \kappa}{n + \kappa}(n\bar{y}^2 + \kappa\mu^2) + S_y + \eta\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 - \frac{(n^2\bar{y}^2 + \kappa^2\mu^2 + 2n\bar{y}\kappa\mu)}{n + \kappa} + \frac{n^2\bar{y}^2 + \kappa n\bar{y}^2 + n\kappa\mu^2 + \kappa^2\mu^2}{n + \kappa} + S_y + \eta\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 + \frac{(-n^2\bar{y}^2 - \kappa^2\mu^2 - 2n\bar{y}\kappa\mu)}{n + \kappa} + \frac{n^2\bar{y}^2 + \kappa n\bar{y}^2 + n\kappa\mu^2 + \kappa^2\mu^2}{n + \kappa} + S_y + \eta\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 + \frac{(-n^2\bar{y}^2 - \kappa^2\mu^2 - 2n\bar{y}\kappa\mu)}{n + \kappa} + \frac{n^2\bar{y}^2 + \kappa n\bar{y}^2 + n\kappa\mu^2 + \kappa^2\mu^2}{n + \kappa} + S_y + \eta\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 + \frac{(-\kappa^2\mu^2 - 2n\bar{y}\kappa\mu + \kappa n\bar{y}^2 + n\kappa\mu^2)}{n + \kappa} + S_y + \eta\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 + \frac{n\kappa(-\kappa\mu^2 - 2\bar{y}\mu + \bar{y}^2 + \mu^2)}{n + \kappa} + S_y + \eta\right)\right\} \\
 &= \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2 + \frac{-n\kappa^2\mu^2 \cdot n\kappa(\bar{y} - \mu)^2}{n + \kappa} + S_y + \eta\right)\right\}
 \end{aligned} \tag{8}$$

So we have

$$p(\theta, \omega | \mathbf{y}) \propto \omega^{(\frac{n+d+1}{2}-1)} \cdot \exp\left\{-\frac{\omega}{2}\left((n + \kappa)\left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2\right)\right\} \cdot \exp\left\{-\frac{\omega}{2}\left(\frac{-n\kappa^2\mu^2 \cdot n\kappa(\bar{y} - \mu)^2}{n + \kappa} + S_y + \eta\right)\right\}$$

which has the form of the normal-gamma

$$p(\theta, \omega | \mathbf{y}) \propto \omega^{(d^*+1)/2-1} \exp\left\{-\omega \cdot \frac{\kappa^*(\theta - \mu^*)^2}{2}\right\} \exp\left\{-\omega \cdot \frac{\eta^*}{2}\right\}$$

From this form of the posterior, the new updated parameters are

- $\mu \rightarrow \mu^* = \frac{n\bar{y} + \kappa\mu}{n + \kappa}$
- $\kappa \rightarrow \kappa^* = n + \kappa$ AND $d \rightarrow d^* = n + d$
- $\eta \rightarrow \eta^* = \frac{-n\kappa^2\mu^2 \cdot n\kappa(\bar{y} - \mu)^2}{n + \kappa} + S_y + \eta$

- (C) From the joint posterior just derived, what is the conditional posterior distribution $p(\theta | \mathbf{y}, \omega)$?
 - you can read it off directly from the joint distribution, since you took care to set up things so that the joint posterior was in the same form as Equation 2.

* normal - gamma (normal sampling model with gamma prior on precision) from Eq 2:

$$p(\theta, \omega) \propto \omega^{\frac{d+1}{2}-1} \cdot \exp\left(-\omega \frac{\kappa(\theta - \mu)^2}{2}\right) \cdot \exp\left(-\omega \cdot \frac{\eta}{2}\right)$$

In general $p(\theta, \omega) \propto p(\omega)p(\theta | \omega)$

and $p(\omega) = \omega^{\frac{d+1}{2}-1} \cdot \left(-\omega \cdot \frac{\eta}{2}\right)$ and $p(\theta | \omega) = \exp\left(-\omega \frac{\kappa(\theta - \mu)^2}{2}\right)$

Thus reading off from B,

$$p(\theta | \mathbf{y}, \omega) \propto \exp\left\{-\frac{\omega}{2} \left((n + \kappa) \left(\theta - \frac{n\bar{y} + \kappa\mu}{n + \kappa}\right)^2\right)\right\}$$

This is the Normal distribution form, $p(\theta | \mathbf{y}, \omega) \sim N\left(\frac{n\bar{y} + \kappa\mu}{n + \kappa}, -\omega(n + \kappa)\right)$

- (D) From the joint posterior calculated in (B), what is the marginal posterior distribution $p(\omega | y)$?
 - Unlike the previous question, where you could just read it off, here you have to integrate over θ . Ignore constants not depending on ω in calculating this integral.

$$\begin{aligned} p(\omega | y) &= \int p(\theta, \omega | y) d\theta \\ &\propto \int \omega^{\frac{d^*+1}{2}-1} \exp\left\{-\omega \cdot \frac{\kappa^*(\theta - \mu^*)^2}{2}\right\} \exp\left\{-\omega \cdot \frac{\eta^*}{2}\right\} d\theta \\ &= \omega^{\frac{d^*+1}{2}-1} \cdot \exp\left\{-\omega \cdot \frac{\eta^*}{2}\right\} \int \exp\left\{-\omega \cdot \frac{\kappa^*(\theta - \mu^*)^2}{2}\right\} d\theta \end{aligned} \quad (9)$$

The integral here is a Normal kernel: $N(\mu^*, \omega\kappa^*)$ and as a probability distribution must integrate to $\frac{1}{c}$, where c is the constant of proportionality for the Normal density, which in this case is $\omega^{1/2}$. Thus $\frac{1}{c} = \frac{1}{\omega^{1/2}} = \omega^{-1/2}$

$$\begin{aligned} &= \omega^{\frac{d^*+1}{2}-1} \cdot \exp\left\{-\omega \cdot \frac{\eta^*}{2}\right\} \cdot \omega^{-1/2} \\ &= \omega^{\frac{d^*}{2}-1} \cdot \exp\left\{-\omega \cdot \frac{\eta^*}{2}\right\} \end{aligned} \quad (10)$$

This is the kernel for the gamma distribution, $\text{Gamma}\left(\frac{d^*}{2}, \frac{\eta^*}{2}\right)$ as seen in Eq 2.

- (E) From (C) and (D), we know that the marginal posterior distribution $p(\theta|\mathbf{y})$ is a gamma mixture of normals. Show that this takes the form of a centered, scaled t distribution:

$$p(\theta) \propto \left(1 + \frac{1}{v} \cdot \frac{(x - m)^2}{s^2}\right)^{-\frac{v+1}{2}}$$

with center m , scale s , and degrees of freedom v (fill in the blank for m , s^2 , and v). express the parameters of this t distribution in terms of the four parameters of the normal-gamma posterior for (θ, ω) . Note: since you've set up the normal-gamma family in this careful conjugate form, this should require no extra work. It's just part (A), except for the prior rather than the posterior.

$$\begin{aligned} p(\theta) &= \int p(\theta, \mathbf{y}) d\mathbf{y} \\ &\propto \int \mathbf{y}^{\frac{d+1}{2}-1} \cdot \exp\left(-\mathbf{y} \frac{\kappa(\theta - \mu)^2}{2}\right) \cdot \exp\left(-\mathbf{y} \cdot \frac{\eta}{2}\right) \\ &= \int \mathbf{y}^{\frac{d+1}{2}-1} \cdot \exp\left(\frac{\mathbf{y} \kappa(\theta - \mu)^2 + -\mathbf{y} \eta}{2}\right) \\ &= \int \mathbf{y}^{(\frac{d+1}{2}-1)} \cdot \exp\left(-\mathbf{y} \left(\frac{\kappa(\theta - \mu)^2 + \eta}{2}\right)\right) \end{aligned} \quad (11)$$

This is the kernel for gamma(a,b) for $a = \frac{d+1}{2}$ and $b = \frac{\kappa(\theta - \mu)^2 + \eta}{2}$

As a probability distribution, it must integrate to $1/c$ where c is a constant, in this case that for the gamma distribution pdf, $c = \frac{b^a}{\Gamma(a)}$ thus $1/c = \frac{\Gamma(a)}{b^a}$ so we can rewrite (1) as:

$$\begin{aligned} &= \Gamma\left(\frac{d+1}{2}\right) \left(\frac{\kappa(\theta - \mu)^2 + \eta}{2}\right)^{-\frac{d+1}{2}} \\ &\propto \left(\frac{\kappa(\theta - \mu)^2 + \eta}{2}\right)^{-\frac{d+1}{2}} \\ &= \left(\frac{\eta}{2} + \frac{\kappa(\theta - \mu)^2}{2}\right)^{-\frac{d+1}{2}} \\ &= \left(\frac{\eta}{2} \cdot \left(1 + \frac{\kappa(\theta - \mu)^2}{\eta}\right)\right)^{-\frac{d+1}{2}} \\ &= \frac{\eta^{-\frac{d+1}{2}}}{2} \cdot \left(1 + \frac{\kappa(\theta - \mu)^2}{\eta}\right)^{-\frac{d+1}{2}} \end{aligned} \quad (12)$$

Since we only care about the value of the equation with respect to θ , we can treat the first part of the equation as constant

$$\begin{aligned} &= \left(1 + \frac{\kappa(\theta - \mu)^2}{\eta}\right)^{-\frac{d+1}{2}} \\ &= \left(1 + \frac{d}{d} \cdot \frac{\kappa(\theta - \mu)^2}{\eta}\right)^{-\frac{d+1}{2}} \\ &= \left(1 + \frac{1}{d} \cdot \frac{(\theta - \mu)^2}{\eta d / \kappa}\right)^{-\frac{d+1}{2}} \end{aligned} \quad (13)$$

This is close to the form of our solution, $p(\theta) \propto \left(1 + \frac{1}{v} \cdot \frac{(x-m)^2}{s^2}\right)^{-\frac{v+1}{2}}$

Thus set $v = d$, $m = \mu$ and set $s^2 = \eta d / \kappa$ to get the centered, scaled t-student form.

- (F) *True or false:* in the limit as the prior parameters κ , d , and η approach zero, the priors $p(\theta)$ and $p(\omega)$ are valid probability distributions.
- *a valid probability distribution must integrate to 1 (or something finite, so that it can be normalized to integrate to 1) over its domain.*
- (G) *True or false:* in the limit as the prior parameters κ , d , and η approach zero, the posteriors $p(\theta|\mathbf{y})$ and $p(\omega|\mathbf{y})$ are valid probability distributions.
- (H) Your result in (F) implies that a Bayesian credible interval for θ takes the form

$$\theta \in m \pm t^* \cdot s,$$

where m and s are the posterior center and scale parameters from (F), and t^* is the appropriate critical value of the t distribution for your coverage level and degrees of freedom (e.g. it would be 1.96 for a 95% interval under the normal distribution).

True or false: In the limit as the prior parameters κ , d , and η approach zero, the Bayesian credible interval for θ becomes identical to the classical (frequentist) confidence interval for θ at the same confidence level.

Appendix A

R code